


P2 – TRAFFIC SIGN CLASSIFIER

SELF DRIVING CAR NANODEGREE



PRATHEERTH PADMAN

Goals

The goals / steps of this project are the following:

- Load the data set (see below for links to the project data set)
- Explore, summarize and visualize the data set
- Design, train and test a model architecture
- Use the model to make predictions on new images
- Analyze the softmax probabilities of the new images
- Summarize the results with a written report

Rubric Points

Here I will consider the [rubric points](#) individually and describe how I addressed each point in my implementation.


Data Set Summary & Exploration:

I used the pandas library to calculate summary statistics of the traffic signs data set:


- The size of training set is - 34,799
- The size of the validation set is – 4,410
- The size of test set is – 12,630
- The shape of a traffic sign image is – (32, 32, 3)
- The number of unique classes/labels in the data set is – 43

Visualization:


1. Training Data



2. Testing Data



3. Validation Data



4. Image Sampling

I looked at 15 random images from the training data to see the quality and type of images we're dealing with.



Design and Test a Model Architecture:

The model architecture is contained in Code Cell 17 in the Jupyter Notebook.

My model is a modification of the famous LeNet model for image classification.

Firstly, I decided to grayscale the given images. Since this is a comparatively simple model, gradient descent will have a simpler time converging because it has to optimize the weights and biases only on a single color channel.

Next, I normalized the image data within a range of -1 to +1. I did it using the formula:

$$X_{\text{normal}} = (X/122.5) - 1$$

Normalizing our image data is important as it keeps our input feature values within a fixed range which helps with the networks learning ability. If input features aren't of a similar scale, there is always a probability that one feature, due to its values, completely overwhelms the others.

To make it "harder" for the network during the training phase, so that it has an easier time during the testing phase, I decided to augment the training dataset.

I augmented the dataset by applying affine transformations (transformations where parallel lines before transformation remain parallel even after transformations) to each image in the training dataset such as rotations, translations and shearing.

The augmentation process can be found in Code Cells 7 and 8.

While I noticed the difference in the frequency of occurrence of each class in the datasets, I decided not to make them equal in order to give the

model a more “real-world” dataset to work with.

Model Description:

My final model consisted of the following layers:

Layer	Description
Input	32x32x1 Grayscale image
Convolution 5x5	1x1 stride, “VALID” padding, outputs 28x28x6
ReLU	Activation Layer
Dropout	Keep Probability – 0.75
Max pooling	2x2 stride, outputs 14x14x6
Convolution 5x5	1x1 stride, “VALID” padding, outputs 10x10x16
ReLU	Activation Layer
Dropout	Keep Probability – 0.75
Max Pooling	2x2 stride, outputs 5x5x16
Flatten	Outputs 400
Fully Connected	Output - 120
ReLU	Activation Layer
Dropout	Keep Probability – 0.75
Fully Connected	Output - 84
ReLU	Activation Layer
Dropout	Keep Probability – 0.75
Fully Connected	Output - 43

Model Training:

The code and the hyper parameters used can be found on Code Cells 19, 20 and 21.

I used an Adam Optimizer as specified in the LeNet Model. The batch size was 256, learning rate was kept at 0.001 and the number of epochs used was 30, which as you can see from the Loss graph below was just enough.

Training Approach:

Initially, I had trained my model using the LeNet model as it stood, only changing small factors in order to properly reflect our dataset. It gave me an expected validation accuracy of ~ 0.89 .

I gray scaled and normalized the images which increased the validation accuracy to ~ 0.91

Including the augmented data in the training set and using dropout layers (which force the network to learn redundant information about the input features while reducing co-adaptability and co-dependence, thus preventing overfitting) are what really caused the accuracy to take off and settle at around $\sim 0.96+$

My final model results were:

- training loss – 0.084
- validation set accuracy - 0.965
- test set accuracy - 0.932

If an iterative approach was chosen:

- **What was the first architecture that was tried and why was it chosen?**

The first architecture that was tried was the LeNet model as it was in order to gauge the performance of the model on our dataset and to set a baseline

- **What were some problems with the initial architecture?**

The model was too simple to reach a validation accuracy of 0.93 +

- **How was the architecture adjusted and why was it adjusted?**

As explained above.

- **Which parameters were tuned? How were they adjusted and why?**

The learning rate was kept at 0.001 as decreasing it would mean increasing the number of epochs to stabilize which would mean increase in the time required to train.

The number of epochs was raised to 30. This represented enough time to properly train the model without compensating model performance for time taken. It also does not overfit as seen from the loss graph.

The probability of keeping a unit in the dropout layer was at 0.75 which represented an ideal tradeoff between optimal training and under-training.

- **Why did you believe it would be relevant to the traffic sign application?**

LeNet model is known to work well on digit classification tasks such as the MNIST model. It would thus also make sense on something like a traffic sign classification application.

- **How does the final model's accuracy on the training, validation and test set provide evidence that the model is working well?**

Firstly, the model can be seen as starting from a high training and validation loss which reduces with each epoch.

We can see that the model does not overfit because of the gap in the validation and training loss.

Lastly, the model can be seen generalizing well as seen from the test set accuracy of 0.932.

Test a Model on New Images:

Here are 10 German traffic signs that I found on the web:



I have the benefit of hindsight here, as while using an earlier trained model with similar test set accuracy, I found that the model suffered with predicting “real world” images which included weird angles, additional imagery and text. The model seemed to do well with properly cropped, “official” traffic sign images.

With this in mind, I expect it to have some difficulty in predicting the 30 km/h speed limit sign as well as the priority road sign.

Model's performance on new images:

Here are the results of the prediction:

Image	Prediction
Right of way at next intersection	Right of way at next intersection
Keep Right	Keep Right
Stop	Yield
30 km/h speed limit	30 km/h speed limit
Bumpy Road	Road narrows on the right
Priority Road	Priority Road
No vehicles	Keep Right
No passing	No passing
No entry	No entry
Children crossing	Children Crossing

The model correctly classifies 7 out of 10 images correctly for an accuracy of 70%.

This obviously is well below the testing and validation set accuracy which might point to overfitting. It could also be down to the fact that we have only taken 10 new images and misclassifying even one of them leads to a 10% reduction in accuracy. The model might do better on a larger dataset.

Softmax Probabilities:

The code for the softmax probabilities is in Code Cell 63.

For the first image, the model is completely sure that this is a right of way sign and the image does contain it. The top five soft max probabilities were

Probability	Prediction
1.0	Right of way at next intersection
0	Priority Road
0	Beware of ice
0	Slippery road
0	Double curve

For the second image:

Probability	Prediction
1.0	Keep right
0	Turn left ahead
0	Go straight or right
0	Slippery road
0	Yield

For the third image:

Probability	Prediction
0.99000001	Yield
0.007	Go straight or right
0.001	Right of way
0.001	Keep Right
0.001	Speed 60 km/h

For the fourth image:

Probability	Prediction
0.8333	Speed 30km/h
0.145	Speed 50km/h
0.021	Speed 20km/h
0	Speed 80 km/h
0	Speed 100 km/h

For the fifth image:

Here, in one of the misclassified image, the model seems to have trouble and it nearly predicted the right sign.

Probability	Prediction
0.563	Road narrows on right
0.3829	Bumpy road
0.053	Bicycles crossing
0.001	Traffic signals
0	Pedestrians

For the sixth image:

Probability	Prediction
0.985	Priority Road
0.009	Yield
0.004	Ahead only
0.001	Speed 60 km/h
0	Turn left ahead

For the seventh image:

Here, even though the second option is correct, the model seems to overwhelmingly believe that the first, wrong prediction is right.

Probability	Prediction
0.995	Keep right
0.003	No vehicles
0	Ahead only
0	End of speed limits
0	Priority road

For the eighth image:

Probability	Prediction
1.0	No passing
0	Priority road
0	No entry
0	No vehicles > 3.5 T
0	End of no passing

For the ninth image:

Probability	Prediction
1.0	No entry
0	Stop
0	Turn left ahead
0	No passing for vehicles > 3.5 T
0	Vehicles over 3.5 T prohibited

For the tenth image:

Probability	Prediction
1.0	Children crossing
0	Dangerous curve to right
0	Right of way
0	Go straight or right
0	Slippery road